R language: a tour

2025-01-14

M1 MIDS/MFA/LOGOS Université Paris Cité Année 2024 Course Homepage Moodle



Objectives

This workbook intends to walk you through basic aspects of the R language and programming environment.

Packages

Base R can do a lot. But the full power of R comes from a fast growing collection of packages.

Packages are first *installed* (that is downloaded from **cran** and copied somewhere on the hard drive), and if needed, *loaded* during a session.

- Installation can usually be performed using command install.packages(). In some circumstances, ad hoc installation commands (often from packages devtools) are needed
- Package pak offers an interesting alternative to base R install.packages()
- Once a package has been installed/downloaded on your hard drive
 - if you want all objects exported by the package to be available in your session, you should load the package, using library() or require() (what's the difference?).
 Technically, this loads the NameSpace defined by the package.
 - if you just want to pick some objects exported from the package, you can use qualified names like package_name::object_name to access the object (function, dataset, ...).

For example, when we write

gapminder <- gapminder::gapminder</pre>

we assign dataframe/tibble gapminder from package gapminder to identifier "gapminder" in global environment .

Function p_load() from pacman (package manager) blends installation and loading: if the package named in the argument of p_load() is not installed (not among the installed.packages()), p_load() attempts to install the package. If installation is successful, the package is loaded.

```
if (! require(pak)){
   install.packages("pak")
}

stopifnot(
   require("tidyverse"),
   require("lobstr"),
   require("ggforce"),
   require("mycflights13"),
   require("patchwork"),
   require("viridis"),
   require("WASS"),
   require("gapminder"),
   require("pryr"),
   require("pryr"),
   require("pak")
)
```

i Optional arguments

A very nice feature of R is that functions from base R as well as from packages have *optional* arguments with sensible *default* values. Look for example at documentation of require() using expression ?require.

Optional settings may concern individual functions or the collection of functions exported by some packages. In the next *chunk*, we reset the default color scales used by graphical functions from ggplot2.

```
opts <- options() # save old options

options(ggplot2.discrete.colour="viridis")
options(ggplot2.continuous.colour="viridis")</pre>
```

You shall not confuse installing (on your hard-drive) and loading (in session) a package.

i Question for Pythonistas

- In **?** what is the analogue of install.packages()?
- In **\rightharpoonup** what is the analogue of require()/library()?

Solution

In **\dial**, you can install a package pck using pip install pck or conda install pck (for example).

In **\(\frac{1}{2} \)**, the analogue of require(pck) could be

```
from pck import *
```

Note that in R, once a package in installed on the hard drive, you do not need to write something like

```
import pck
```

to be able to use objects exported by pck using qualified names (like pck.ze_object), you just need to use R qualified names:

```
pck::ze_object
```

Numerical (atomic) vectors

Numerical (atomic) vectors form the most primitive type of R.

Vector creation and assignment

The next three lines create three numerical atomic vectors.

In IDE Rstudio, have a look at the environment pane on the right before running the chunk, and after.

Use ls() to investigate the *environment* before and after the execution of the three assignments.

```
1 ls()
2 x <- c(1, 2, 12)
3 y <- 5:7
4 z <- 10:1
5 x; y; z
6 ls()
```

i Question

Solution

x ; y ; z

- What are the identifiers known in the global environment before execution of lines 2-4?
- What are the identifiers known in the global environment after execution of lines 2-4?
- Which objects are attached to identifiers x, y, and z?

```
ls()
[1] "opts"

x <- c(1, 2, 12)

y <- 5:7

z <- 10:1
```

```
[1] 1 2 12
[1] 5 6 7
[1] 10 9 8 7 6 5 4 3 2 1

ls()
[1] "opts" "x" "y" "z"

The chunks adds three identifiers x,y,z to the global environment. Identifiers are bound to R objects which turn out to be numerical vectors.
```

Question

What does the next chunk?

```
ls()
w <- y
ls()
```

Solution

The chunk inserts a new identifier w in the global environment. This identifier is associated with the same object as y.

i Question

- Is the content of object denoted by y copied to a new object bound to w?
- Interpret the result of w == y.

The address associated with y has changed!

• Interpret the result of identical(w,y) (use help("identical") if needed).

```
w == y
identical(w,y)
```

Solution

Package lobstr lets us explore low-level aspects of R (and much more). Function lobstr::obj_addr() returns the address of the object denoted by the argument.

```
lobstr::obj_addr() returns the address of the object denoted by the argument.
lobstr::obj_addr(w)
[1] "0x5a0e3adf5938"
lobstr::obj_addr(y)
[1] "0x5a0e3adf5938"
Now, if we modify either y or w
y <- y + 1
identical(y, w)
[1] FALSE
c(lobstr::obj_addr(w), lobstr::obj_addr(y))
[1] "0x5a0e3adf5938" "0x5a0e355bb678"</pre>
```



The meaning of assignment in R differs from its counterpart in Python. In Python, assignment is shallow. In R, assignment creates a new identifier bound to the same object as the right-hand side of the assignment. If either side of the assignment is modified, it is copied to a new object before modification. This is called *copy-on-modify*.

Indexation, slicing, modification

Slicing a vector can be done in two ways:

- providing a vector of indices to be selected. Indices need not be consecutive.
- providing a Boolean mask, that is a logical vector to select a set of positions.

```
x \leftarrow c(1, 2, 12); y \leftarrow 5:7; z \leftarrow 10:1
```

i Question

Explain the next lines

```
z[1]  # slice of length 1
z[0]  # What did you expect?
z[x]  # slice of length ??? index error ?
z[y]
z[x %% 2]  # what happens with x[0] ?
z[0 == (x %% 2)]  # masking
z[c(2, 1, 1)]
```

Solution

- Indices start at 1 (not like in C, Java, or Python)
- \mathbf{z} [0] does not return an Error message. It returns an empty vector with the same basetype as \mathbf{x}
- z[x] returns a vector made of z[x[1]], z[x[2]] and z[x[3]]==z[12]. Note again that z[12] does not raise an exception. It is simply not available (NA).
- x % 2 returns 1 0 0 as % stands for mod. z[x % 2] returns the same thing as z[1]
- c() stands for combine, or concatenate.

Question

If the length of mask and and the length of the sliced vector do not coincide, what happens?

Solution

No error is signaled, the returned sequence is as long as the number of truthies in the mask.

Out of bound truthies show up as NA

```
z[rep(c(TRUE, FALSE), 6)]
[1] 10 8 6 4 2 NA
```

i

A scalar is just a vector of length 1!

```
class(z)
[1] "integer"

class(z[1])
[1] "integer"

class(z[c(2,1)])
[1] "integer"
```

i Question

Explain the next lines

```
y[2:3] <- z[2:3]

y == z[-10]

z[-11]
```

Solution

We can assign a slice of a vector to a slice of identical size of another vector. What is the result of z[-11], z[-c(11:7)]?

Question

Explain the next line

```
z[-(1:5)]
```

Solution

We pick all positions in z but the ones in 1:5, that is 6, 7, 8, 9, 10

i Question

How would you select the last element from a vector (say z)?

Solution

```
z[length(z)]
```

[1] 1

6 10

Q is not **♦** (reminder)!

i Question

Reverse the entries of a vector. Find two ways to do that.

```
Solution

z[seq(length(z), 1, by=-1)]

[1] 1 2 3 4 5 6 7 8 9 10

z[length(z):1]

[1] 1 2 3 4 5 6 7 8 9 10

rev(z) # the simplest way, once you know rev()

[1] 1 2 3 4 5 6 7 8 9 10
```

In statistics, machine learning, we are often faced with the task of building grids of regularly spaced elements (these elements can be numeric or not). R offers a collection of tools to perform this. The most basic tool is rep().

Question

- Repeat a vector 2 times
- Repeat each element of a vector twice

```
Solution

w <- c(1, 7, 9)
rep(w, 2)

[1] 1 7 9 1 7 9

rep(w, rep(2, length(w)))

[1] 1 1 7 7 9 9

Now, we can try something more fancy.

rep(w, 1:3)

[1] 1 7 7 9 9 9

What are the requirements on the second (times) argument?
```

Let us remove objects from the global environment.

```
rm(w, x, y ,z)
```

Numbers

So far, we told about numeric vectors. Numeric vectors are vectors of floating point numbers. R distinguishes several kinds of numbers.

- Integers
- Floating point numbers (double)

To check whether a vector is made of numeric or of integer, use is.numeric() or is.integer(). Use as.integer, as.numeric() to enforce type conversion.

```
i Question
Explain the outcome of the next chunks

class(113L); class(113); class(113L + 113); class(2 * 113L); class(pi); as.intege
[1] "integer"
[1] "numeric"
[1] "numeric"
[1] "numeric"
[1] "numeric"
[1] 3

floor(pi); class(floor(pi)) # mind the floor
[1] 3
[1] "numeric"
```

Integer arithmetic

```
29L * 31L ; 899L %/% 32L ; 899L %% 30L

[1] 899

[1] 28

[1] 29
```

```
R integers are not the natural numbers from Mathematics
R numerics are not the real numbers from Mathematics
.Machine$double.eps
[1] 2.220446e-16
.Machine$double.xmax
[1] 1.797693e+308
.Machine$sizeof.longlong
[1] 8
u <- double(19L)
v <- numeric(5L)</pre>
w <- integer(7L)
lapply(list(u, v, w), typeof)
[[1]]
[1] "double"
[[2]]
[1] "double"
[[3]]
[1] "integer"
length(c(u, v, w))
[1] 31
typeof(c(u, v, w))
[1] "double"
```

R is (sometimes) able to make sensible use of Infinite.

```
log(0)
[1] -Inf
log(Inf)
[1] Inf
1/0
[1] Inf
0/0
[1] NaN
\max(c(0/0,1,10))
[1] NaN
\max(c(NA,1,10))
[1] NA
max(c(-Inf,1,10))
[1] 10
```

```
is.finite(c(-Inf,1,10))
[1] FALSE TRUE TRUE
is.na(c(NA,1,10))
[1] TRUE FALSE FALSE
is.nan(c(NaN,1,10))
[1] TRUE FALSE FALSE
Computing with vectors
Summing, scalar multiplication
x <- 1:3
y <- 9:7
sum(x); prod(x)
[1] 6
[1] 6
z <- cumsum(1:3)
w <- cumprod(3:5)</pre>
x + y
[1] 10 10 10
x + z
[1] 2 5 9
2 * w
[1] 6 24 120
2 + w
[1] 5 14 62
w / 2
[1] 1.5 6.0 30.0
  Question
    How would you compute a factorial?
  Solution
  n <- 10
  cumprod(1:n)
                     2
   [1]
                             6
                                    24
                                           120
                                                   720
                                                           5040
                                                                  40320 362880
             1
  [10] 3628800
```

i Question

Approximate $\sum_{n=1}^{\infty} 1/n^2$ within 10^{-3} ?

Solution

$$\sum_{n>N} \frac{1}{n^2} < \sum_{n>N} \frac{1}{n(n-1)} = \sum_{n>N} \left(\frac{1}{n-1} - \frac{1}{n} \right) = \frac{1}{N}$$

So we may pick N = 1000.

```
sum(x*y) # inner product
```

[1] 46

```
prod(1:5) # factorial(n) as prod(1:n)
```

[1] 120

```
N \leftarrow 1000L

sum(1/((1:N)^2)); pi^2/6 # grand truth
```

- [1] 1.643935
- [1] 1.644934

```
(pi^2/6 - sum(1/((1:N)^2))) < 1e-3
```

[1] TRUE

```
# N <- 999L
# (pi^2/6 - sum(1/((1:N)^2))) < 1e-3
```

i Question

How would you compute the inner product between two (atomic numeric) vectors?

Solution

Inner product between two vectors can be computed as a matrix product between a row vector and a column vector using %*%. Is this a good idea?

```
matrix(w, ncol=3) %*% matrix(y, nrow=3) == sum(w * y)
        [,1]
[1,] TRUE
```

- What we have called vectors so far are indeed atomic vectors.
 - Read Chapter on Vectors in R advanced Programming
 - Keep an eye on package vctrs for getting insights into the R vectors.

Numerical matrices

R offers a matrix class.

```
A <- matrix(1:50, nrow=5)
A
```

```
[,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [1,] 1 6 11 16 21 26 31 36 41 46
```

[2,]	2	7	12	17	22	27	32	37	42	47
[3,]	3	8	13	18	23	28	33	38	43	48
[4,]	4	9	14	19	24	29	34	39	44	49
[5,]	5	10	15	20	25	30	35	40	45	50

class(A)

[1] "matrix" "array"

i Question

From the evaluation of the preceding chunk, can you guess whether it is easier the traverse a matrix in row-first order or in column-first order?

Solution

Default traversal seems to proceed columnwise.

Creation, transposition and reshaping

A vector can be turned into a column matrix.

```
v <- as.matrix(1:5)
v</pre>
```

[,1] [1,] 1 [2,] 2

[3,] 3

[4,] 4 [5,] 5

A matrix can be transposed

t(v) # transpose

```
[,1] [,2] [,3] [,4] [,5]
[1,] 1 2 3 4 5
cat(dim(v), ' ', dim(t(v)), '\n')
```

5 1 1 5

A <- matrix(1, nrow=5, ncol=2); A

[1,1] [,2]
[1,] 1 1
[2,] 1 1
[3,] 1 1
[4,] 1 1
[5,] 1 1

Question

lobstr::mem_used() allows us to keep track of the amount of memory used by
our R session. lobstr::obj_size() tells us the amount of memory used by the
representation of an object.

Comment the next chunk

```
m1 <-lobstr::mem_used()
A <- matrix(rnorm(100000L), nrow=1000L)
m2 <- lobstr::mem_used()
lobstr::obj_size(A)

800.22 kB
B <- t(A)
lobstr::obj_size(B)

800.22 kB
m3 <- lobstr::mem_used()
m2-m1; m3-m2

802.44 kB
1.09 MB</pre>
```

i Question

- Is there a difference between the next two assignments?
- How would you assign value to all entries of a matrix?

```
A <- matrix(rnorm(16), nrow=4)
A[] \leftarrow 0 ; A
     [,1] [,2] [,3] [,4]
[1,]
       0 0
                 0
             0
                  0
                       0
[2,]
        0
          0
[3,]
       0
                  0
                       0
[4,]
                  0
                       0
       0
             0
A < -0; A
[1] 0
```

Solution

There is!

The first assignment assigns 0 to every entry in A.

The second assignment binds 0 to name A

Question

```
What is the final shape of A?
```

We can easily generate diagonal matrices and constant matrices.

```
diag(1, 3) # building identity matrix
     [,1] [,2] [,3]
[1,]
              0
[2,]
        0
              1
                   0
[3,]
        0
              0
                   1
matrix(0, 3, 3) # building null matrix
     [,1] [,2] [,3]
[1,]
        0
              0
[2,]
                   0
        0
              0
[3,]
              0
                   0
    Question
     Is there any difference between the next two assignments?
     B <- A[]
     B ; A
     [1] 0
     [1] 0
    lobstr::obj_addr(B) ; lobstr::obj_addr(A)
     [1] "0x5a0e3a1bc610"
     [1] "0x5a0e3a942f38"
     B <- A
     lobstr::obj_addr(B) ; lobstr::obj_addr(A)
     [1] "0x5a0e3a942f38"
     [1] "0x5a0e3a942f38"
```

Indexation, slicing, modification

Indexation consists in getting one item from a vector/list/matrix/array/dataframe.

Slicing and subsetting consists in picking a substructure:

- subsetting a vector returns a vector
- subsetting a list returns a list
- subsetting a matrix/array returns a matrix/array (beware of implicit simplifications and dimension dropping)
- subsetting a dataframe returns a dataframe or a vector (again, beware of implicit simplifications).

i Question

Explain the next results

```
A <- matrix(1, nrow=5, ncol=2)

dim(A[sample(5, 3), -1])
dim(A[sample(5, 3), 1])
length(A[sample(5, 3), 1])
is.vector(A[sample(5, 3), 1])
A[10:15]
A[60]
dim(A[])</pre>
```

```
NULL
NULL
[1] 3
[1] TRUE
[1] 1 NA NA NA NA NA
[1] NA
[1] 5 2
```

i Question

How would you create a fresh copy of a matrix?

Solution

```
A <- matrix(rnorm(10), ncol=2L)
B <- matrix(0, nrow=5L, ncol=2L)

B[] <- A
all(B==A) ; identical(A, B) ; lobstr::obj_addrs(list(A, B))

[1] TRUE
[1] TRUE
[1] "0x5a0e39a04ea8" "0x5a0e39a056e8"</pre>
```

Computing with matrices

* versus %*% %*% stands for matrix multiplication. In order to use it, the two matrices should have conformant dimensions.

```
[,1] [,2]
[1,] -0.6941609 -0.3425532
```

There are a variety of reasonable products around. Some of them are available in R.

Question

How would you compute the Hilbert-Schmidt inner product between two matrices?

$$\langle A, B \rangle_{\mathrm{HS}} = \mathrm{Trace}(A \times B^{\top})$$

Solution

In R, trace() does not return the trace of a matrix! Function is used for debugging. Just remember that the trace of a matrix is the sum of its diagonal elements.

```
A <- matrix(runif(6), 2, 3)

B <- matrix(runif(6), 2, 3)

foo <- sum(diag(A %*% t(B)))

bar <- sum(A * B)

foo ; bar

[1] 0.583757

[1] 0.583757

Are you surprised?
```

i Question

How can you invert a square (invertible) matrix?

Use solve(A) which is a shorthand for solve(A, diag(1, nrow(3))).

Logicals

- R has constants TRUE and FALSE.
- Numbers can be coerced to logicals.

Question

- Which numbers are truthies? falsies?
- What is the value (if any) of ! pi & TRUE ?
- What is the meaning of all()?
- What is the meaning of any()?
- Recall De Morgan's laws. Check them with R.
- Is | denoting an inclusive or an exclusive OR?

```
Solution
  w <- c(TRUE, FALSE, FALSE)
  sum(w)
  [1] 1
  any(w)
  [1] TRUE
  all(w)
  [1] FALSE
  ! w
  [1] FALSE TRUE TRUE
  TRUE & FALSE
  [1] FALSE
  TRUE | FALSE
  [1] TRUE
  TRUE | TRUE
  [1] TRUE
```

Handling three-valued logic

```
TRUE & (1> (0/0))
(1> (0/0)) | TRUE
(1> (0/0)) | FALSE
TRUE || (1> (0/0))
TRUE | (1> (0/0))
TRUE | (1> (0/0))
TRUE || stopifnot(4<3)
# TRUE | stopifnot(4<3)
# FALSE && stopifnot(4<3)
# FALSE & stopifnot(4<3)</pre>
```

Solution TRUE & (1> (0/0)) [1] NA (1> (0/0)) | TRUE [1] TRUE (1> (0/0)) | FALSE [1] NA TRUE || (1> (0/0)) [1] TRUE TRUE | (1>(0/0))[1] TRUE TRUE || stopifnot(4<3) [1] TRUE # TRUE | stopifnot(4<3)</pre> FALSE && stopifnot(4<3)</pre> [1] FALSE # FALSE & stopifnot(4<3)</pre>

i Question

What is the difference between logical operators | | and | ?

Solution

 $| \ |$ is lazy. It does not evaluate its second argument if the first one evaluates to TRUE. && is also lazy.

1

Remark: favor &, | over &&, | |.

all and any

Look at the definition of all and any.

i Question

- How would you check that a square matrix is symmetric?
- How would you check that a matrix is diagonal?

Solution

A square matrix is symmetric iff it is equal to its transpose. Recall that t(A) denotes the transpose of matrix A.

```
A <- matrix(rnorm(9), nrow=3, ncol=3) # a.s. non-symmetric all(A == t(A))
```

```
[1] FALSE

A <- A %*% t(A) # build a symmetric matrix, A + t(A) would work also
all(A == t(A))
[1] TRUE
A == t(A) returns a matrix a logical matrix, whose entries are all TRUE iff A is symmetric.
all() works for matrices as well as for vectors. This is sensible as matrices can be considered as vectors with some additional structure.</pre>
```

Lists

While an instance of an atomic vector contains objects of the same type/class, an instance of list may contain objects of widely different types.

```
Check an explain the output of the next chunk

p <- c(2, 7, 8)
q <- c("A", "B", "C")
x <- list(p, q)
x[2]
x
length(x)
rlang::is_vector(x)
rlang::is_atomic(x)
y <- c(p, q)
y
length(y)
rlang::is_atomic(y)
rlang::is_atomic(y)
rlang::is_list(y)</pre>
```

```
[[1]]
[1] "A" "B" "C"
[[1]]
[1] 2 7 8

[[2]]
[1] "A" "B" "C"
[1] 2
[1] TRUE
[1] FALSE
[1] "2" "7" "8" "A" "B" "C"
[1] 6
[1] TRUE
[1] FALSE
[1] TRUE
[1] TRUE
```

```
### Question
#| label: enigma-list-2
#| eval: true
lobstr::tree(x)

true
lobstr::tree(x)

true
lobstr::tree(y)
```

<chr [6]>"2", "7", "8", "A", "B", "C"

i Question

- How would you build a list made of p, q, and x?
- What is x[2] made of?
- How does it compare with x[[2]]?

```
Solution
nl \leftarrow list(p=p, q=q, x=x)
nl
$р
[1] 2 7 8
$q
[1] "A" "B" "C"
$x
$x[[1]]
[1] 2 7 8
$x[[2]]
[1] "A" "B" "C"
Note that we have defined a named list. Each = expression, binds the string on the
left-hand side to the object on the right-hand side. List elements can be extracted in
defferent ways.
names(nl)
[1] "p" "q" "x"
nl$q
[1] "A" "B" "C"
nl[["q"]]
[1] "A" "B" "C"
nl[[2]]
[1] "A" "B" "C"
```

Question

Read and understand the next expressions.

```
is_atomic(p); is_atomic(p[2]); is_atomic(p[[2]])
is_list(q); is_atomic(q)
is_list(x); is_atomic(x); class(x)

class(x[2]); class(x[[2]])
length(x[2]); length(x[[2]])

identical(q, x[[2]]); identical(q, x[2])

obj_addr(q); obj_addr(x[[2]]); obj_addr(x[2])
ref(x)
obj_addrs(x)
identical(x[2],x[[2]])
```

i Functions is_atomic(), is_list(), ..., obj_addr() are from packages rlang and lobstr. See https://rlang.r-lib.org and https://lobstr.r-lib.org

Solution

p and a are atomic vectors with different base types. They are not lists. A list like \boldsymbol{x} is not an atomic vector.

Inspection of object addresses shows that when building x from objects p and q, objects bound to "p" and "q" are not copied.

Note that x[[2]] and x[2] are different objects, the former is one element list, the second is an atomic vector.

```
lobstr::ref(x[2])
  [1:0x5a0e3a151600] <list>
  [2:0x5a0e32857858] <chr>
lobstr::obj_addr(x[[2]])
[1] "0x5a0e32857858"
```

Question

How would you replace "A" in x with "K"?

Solution

```
x[[2]][1] <- "K"

x
[[1]]
[1] 2 7 8

[[2]]
[1] "K" "B" "C"
```

```
lobstr::ref(x)
  [1:0x5a0e3a422d38] <list>
  [2:0x5a0e32856138] <dbl>
  [3:0x5a0e3a58cb78] <chr>

w <- c(2, 7, 8)
v <- c("A", "B", "C")
x <- list(w, v)</pre>
```

Read Chapter on Lists in R advanced Programming

Lookup tables (aka dictionaries) using named vectors

A lookup table maps strings to values. It can be implemented using named vectors. If we want to map: "seine" to "75", "loire" to "42", "rhone" to "69", "savoie" to "73" we can proceed in the following way:

```
codes <- c(75L, 42L, 69L, 73L)
names(codes) <- c("seine", "loire", "rhône", "savoie")

codes["rhône"]; codes["aube"]

rhône
   69
<NA>
NA
```

i Question

What is the class of codes?

```
names(codes)
[1] "seine" "loire" "rhône" "savoie"

class(codes); class(names(codes))
[1] "integer"
[1] "character"

is_atomic(codes); is_character(codes) ; is_integer(codes)

[1] TRUE
[1] FALSE
[1] TRUE
```

Question

Capitalize the names used by codes

Package stringr offers a function str_to_title() that could be of interest.

```
Solution

names(codes) <- stringr::str_to_title(names(codes))
codes

Seine Loire Rhône Savoie
75 42 69 73
```

Factors

Factors exist in Base R. They play a very important role. Qualitative/Categorical variables are implemented as Factors.

Meta-package tidyverse offers a package dedicated to factor engineering: forcats.

- [1] TRUE
- [1] TRUE

i Question

g1 g2 g3

yraw takes few values. It makes sense to make it a factor. How does it change the behavior of *generic* function summary?

fyraw <- as.factor(yraw) levels(fyraw) [1] "g1" "g2" "g3" summary(fyraw)</pre>

Load the (celebrated) iris dataset, and inspect variable Species

```
data(iris)
species <- iris$Species
levels(species)</pre>
```

[1] "setosa" "versicolor" "virginica"

summary(species)

```
setosa versicolor virginica
50 50 50
```

i Question

We may want to collapse virginica and versicolor into a single level called versinica

forcats offer a function fct_collapse().

```
Solution

col_species <- forcats::fct_collapse(
   species,
   versinica = c("versicolor", "virginica")
)

summary(col_species)

setosa versinica
   50   100</pre>
```

Factors are used to represent *categorical* variables.

Question

- Load the whiteside data from package MASS.
- Have a glimpse.
- Assign column Insul to y

```
Solution
```

```
whiteside <- MASS::whiteside # importing the whiteside data
# ?whiteside # what are the whiteside data about?

tibble::glimpse(whiteside)

Rows: 56
Columns: 3
$ Insul <fct> Before, Before, Before, Before, Before, Before, Before, Before, Femp <dbl> -0.8, -0.7, 0.4, 2.5, 2.9, 3.2, 3.6, 3.9, 4.2, 4.3, 5.4, 6.0, 6.~
$ Gas <dbl> 7.2, 6.9, 6.4, 6.0, 5.8, 5.8, 5.6, 4.7, 5.8, 5.2, 4.9, 4.9, 4.3,~
y <- whiteside$Insul # picking a factor column</pre>
```

Question

- What is the class of y?
- Is y a vector
- Is y ordered? What does ordered mean here?
- What are the levels of y? How many levels has y?
- Can you slice y?
- What are the binary representations of the different levels of y?

```
Solution
is.factor(y); is.vector(y); is.ordered(y)
[1] TRUE
[1] FALSE
[1] FALSE
class(y)
[1] "factor"
levels(y)
[1] "Before" "After"
nlevels(y)
[1] 2
y[1:10] # yes we can
 [1] Before Before Before Before Before Before Before Before Before
Levels: Before After
pryr::bits(y[31]) # looks like the two levels are represented by integers
[1] "00000000 00000000 00000000 00000010"
```

i Question

Summarize factor y

```
Solution

summary(y) # counts

Before After
    26    30

table(y) # one-way contingency table

y

Before After
    26    30

table(y)/sum(table(y))*100 # one-way contingency table as percentages

y

Before After
    46.42857 53.57143
```

```
table(y) |>
  knitr::kable(col.names = c("Insulation", "Frequency"),
                caption = "Whiteside data") # Pb encoding sur machine windows
                           Table 2: Whiteside data
                            Insulation
                                      Frequency
                            Before
                                              26
                            After
                                              30
forcats::fct_count(y) |>
  knitr::kable(col.names = c("Insulation", "Frequency"),
                caption = "Whiteside data")
                           Table 3: Whiteside data
                            Insulation
                                       Frequency
                            Before
                                              26
                            After
                                              30
```

Factors nuts and bolts

When coercing a vector (integer, character, ...) to a factor, use forcats::as_factor() rather than base R as.factor().



Useful function to make nice barplots when constructing barplots.

Recall that when you want to display counts for a univariate categorical sample, you use a barplot. It is often desirable to rank the levels according to the displayed statistics (usually a count).

This can be done in a seamless way using functions like forcats::fct_infreq().

```
forcats::fct_count(y, prop = TRUE)
```

```
# A tibble: 2 x 3
            n
  <fct> <int> <dbl>
1 Before
            26 0.464
2 After
            30 0.536
z <- sample(y, length(y), replace = TRUE) # permutation of whiteside$Insul
sort(forcats::fct_infreq(z))
                                   # first level is most frequent one
```

- [1] Before Before Before Before Before Before Before Before Before
- [11] Before Befo
- [21] Before Before Before Before Before Before Before After After
- [31] After After After After After After After After After
- [41] After After After After After After After After After After
- [51] After After After After After

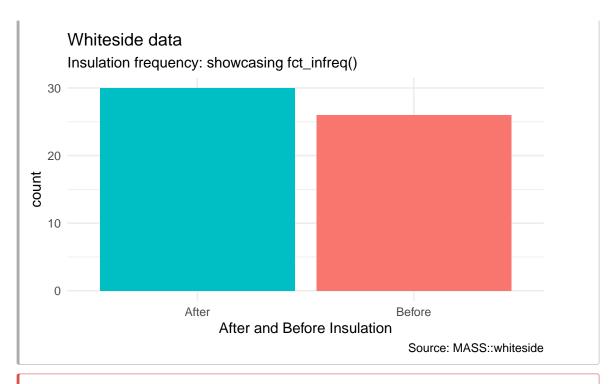
Levels: Before After

i Question

Make z ordered with level After preceding Before. Does ordering impact the behavior of forcats::fct_count()?

```
Solution
forcats::fct_count(z)
# A tibble: 2 x 2
 f
  <fct> <int>
1 Before
           28
2 After
           28
forcats::fct_count(factor(z, ordered=TRUE, levels=c("After", "Before")))
# A tibble: 2 x 2
 f
  <ord> <int>
1 After
           28
2 Before
           28
forcats::fct_count(forcats::fct_infreq(z))
# A tibble: 2 x 2
  <fct> <int>
1 Before
           28
2 After
           28
```

```
whiteside |>
  ggplot2::ggplot() +
  ggplot2::aes(x=forcats::fct_infreq(Insul), fill=Insul) +
  ggplot2::geom_bar() +
  ggplot2::xlab("After and Before Insulation") +
  ggplot2::theme_minimal() +
  ggplot2::theme(legend.position="None") +
  labs(title="Whiteside data",
      subtitle="Insulation frequency: showcasing fct_infreq()",
      caption="Source: MASS::whiteside")
```



Read Chapter on Factors in R for Data Science

Dataframes, tibbles and data.tables

A dataframe is a list of vectors with equal lengths. This is the way R represents and manipulates multivariate samples.

Any software geared at data science supports some kind of dataframe

- Python Pandas
- · Python Dask
- Spark
- ...

The iris dataset is the "Hello world!" of dataframes.

```
data(iris)
iris |>
  glimpse()
```

```
Rows: 150
Columns: 5
$ Sepal.Length <dbl> 5.1, 4.9, 4.7, 4.6, 5.0, 5.4, 4.6, 5.0, 4.4, 4.9, 5.4, 4.~
$ Sepal.Width <dbl> 3.5, 3.0, 3.2, 3.1, 3.6, 3.9, 3.4, 3.4, 2.9, 3.1, 3.7, 3.~
$ Petal.Length <dbl> 1.4, 1.4, 1.3, 1.5, 1.4, 1.7, 1.4, 1.5, 1.4, 1.5, 1.5, 1.5, 1.~
$ Petal.Width <dbl> 0.2, 0.2, 0.2, 0.2, 0.2, 0.4, 0.3, 0.2, 0.2, 0.1, 0.2, 0.~
$ Species <fct> setosa, setosa
```

A matrix can be transformed into a data.frame

```
A <- matrix(rnorm(10), ncol=2)
data.frame(A)
```

```
X1 X2
1 -0.61518158 1.5136984
```

```
2 -0.06755219 0.5215982
3 1.71754357 1.2732472
4 1.05261998 -0.4915186
5 0.06402622 0.4883434
```

There are several flavors of dataframes in R: tibble and data.table are modern variants of data.frame.

```
t <- tibble::tibble(
    x=1:3,
    a=letters[11:13],
    d=Sys.Date() + 1:3)
head(t)</pre>
```

glimpse(t)

```
Rows: 3
Columns: 3
$ x <int> 1, 2, 3
$ a <chr> "k", "l", "m"
$ d <date> 2025-01-15, 2025-01-16, 2025-01-17
```

ref(t)

```
[1:0x5a0e3a2ac808] <tibble[,3]>
x = [2:0x5a0e38adae30] <int>
a = [3:0x5a0e3a372168] <chr>
d = [4:0x5a0e3a315078] <date>
```

Read Chapter on data frames and tibbles in Advanced R

i Question

Perform a random permutation of the columns of a data.frame/tibble.

Function sample() from base R is very convenient

Solution

nycflights data

Wrestling with tables is part of the data scientist job. Out of the box data are often messy. In order to perform useful data analysis, we need *tidy* data. The notion of tidy data was elaborated during the last decade by experienced data scientists.

You may benefit from looking at the following online documents.

Tidy data in R for Data Science

Introduction to Table manipulation in R for Data Science in R.

More data of that kind is available following guidelines from https://github.com/hadley/nycflights13

In this exercise, you are advised to use functions from dplyr.

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges.

```
data <- nycflights13::flights
```

Question

- Have a glimpse at the data.
- What is the class of object data?
- What kind of object is data?

Hint: use class(), is.data.frame() tibble::is_tibble()

```
Solution
#| label: flight_glimpse
#| eval: true
data |> glimpse()
Rows: 336,776
Columns: 19
$ year
              <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2
$ month
              $ day
              $ dep_time
              <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 558, ~
$ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 600, ~
              <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2, -1~
$ dep_delay
$ arr_time
              <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, $49,~
$ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 851,~
```

```
<dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7, -1~
$ arr_delay
                 <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6", "~
$ carrier
$ flight
                 <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301, 4~
                 <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N394~
$ tailnum
                 <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LGA",~
$ origin
                 <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IAD",~
$ dest
$ air_time
                 <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149, 1~
$ distance
                 <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 733, ~
$ hour
                 <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6, 6
$ minute
                 <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 0, 59, 0~
$ time_hour
                 <dttm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-01 0~
class(data)
[1] "tbl_df"
                 "tbl"
                              "data.frame"
is.data.frame(data)
[1] TRUE
is_tibble(data)
[1] TRUE
```

i Question

Extract the name and the type of each column.

```
Solution
colnames(data)
                             # name of columns
 [1] "year"
                       "month"
                                         "day"
                                                           "dep_time"
 [5] "sched_dep_time"
                       "dep_delay"
                                         "arr_time"
                                                           "sched_arr_time"
                                                           "tailnum"
 [9] "arr_delay"
                       "carrier"
                                         "flight"
                       "dest"
                                         "air_time"
                                                           "distance"
[13] "origin"
[17] "hour"
                       "minute"
                                         "time_hour"
                             # old school R, a dataframe is a list
sapply(data, class)
$year
[1] "integer"
$month
[1] "integer"
$day
[1] "integer"
$dep_time
[1] "integer"
$sched_dep_time
[1] "integer"
$dep_delay
[1] "numeric"
```

```
$arr_time
[1] "integer"
$sched_arr_time
[1] "integer"
$arr_delay
[1] "numeric"
$carrier
[1] "character"
$flight
[1] "integer"
$tailnum
[1] "character"
$origin
[1] "character"
$dest
[1] "character"
$air_time
[1] "numeric"
$distance
[1] "numeric"
$hour
[1] "numeric"
$minute
[1] "numeric"
$time_hour
[1] "POSIXct" "POSIXt"
lapply(data, class)
                             # old school R, a dataframe is a list
$year
[1] "integer"
$month
[1] "integer"
$day
[1] "integer"
$dep_time
[1] "integer"
```

```
$sched_dep_time
[1] "integer"
$dep_delay
[1] "numeric"
$arr_time
[1] "integer"
$sched_arr_time
[1] "integer"
$arr_delay
[1] "numeric"
$carrier
[1] "character"
$flight
[1] "integer"
$tailnum
[1] "character"
$origin
[1] "character"
$dest
[1] "character"
$air_time
[1] "numeric"
$distance
[1] "numeric"
$hour
[1] "numeric"
$minute
[1] "numeric"
$time_hour
[1] "POSIXct" "POSIXt"
purrr::map(data, class)
                                    # tidyverse way
$year
[1] "integer"
$month
[1] "integer"
```

```
$day
[1] "integer"
$dep_time
[1] "integer"
$sched_dep_time
[1] "integer"
$dep_delay
[1] "numeric"
$arr_time
[1] "integer"
$sched_arr_time
[1] "integer"
$arr_delay
[1] "numeric"
$carrier
[1] "character"
$flight
[1] "integer"
$tailnum
[1] "character"
$origin
[1] "character"
$dest
[1] "character"
$air_time
[1] "numeric"
$distance
[1] "numeric"
$hour
[1] "numeric"
$minute
[1] "numeric"
$time_hour
[1] "POSIXct" "POSIXt"
```

Compute the mean of the numerical columns

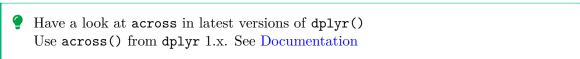
Base R has plenty of functions that perform statistical computations on univariate samples. Look at the documentation of mean (just type ?mean). For a while, leave aside the optional arguments.

In database parlance, we are performing aggregation

```
mean(data$dep_delay)
[1] NA
# mean(data[["dep_delay"]])
```

If we want the mean of all numerical columns, we need to project the data frame on numerical columns.

A verb of the summarize family can be useful.



```
Solution
data |>
  dplyr::select(where(is.numeric)) |> # projecting on numerical columns
  purrr::map(mean)
                                         # applying the treatment to each column
$year
[1] 2013
$month
[1] 6.54851
$day
[1] 15.71079
$dep_time
[1] NA
$sched_dep_time
[1] 1344.255
$dep_delay
[1] NA
$arr_time
[1] NA
$sched_arr_time
[1] 1536.38
$arr_delay
[1] NA
```

```
$flight
[1] 1971.924
$air_time
[1] NA
$distance
[1] 1039.913
$hour
[1] 13.18025
$minute
[1] 26.2301
  dplyr::select(where(is.numeric)) |> # projecting on numerical columns
 dplyr::summarise(across(everything(), mean, na.rm=T))
# A tibble: 1 x 14
                day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  year month
  <dbl> <dbl> <dbl>
                       <dbl>
                                      <dbl>
                                                 <dbl>
                                                          <dbl>
                                                                         <dbl>
                                                          1502.
                                                                         1536.
1 2013 6.55 15.7
                       1349.
                                      1344.
                                                  12.6
# i 6 more variables: arr_delay <dbl>, flight <dbl>, air_time <dbl>,
    distance <dbl>, hour <dbl>, minute <dbl>
data |>
 dplyr::summarise(across(where(is.numeric), mean))
# A tibble: 1 x 14
               day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  year month
  <dbl> <dbl> <dbl>
                       <dbl>
                                      <dbl>
                                                 <dbl>
                                                          <dbl>
1 2013 6.55 15.7
                                                                         1536.
                          NA
                                      1344.
                                                    NA
                                                             NA
# i 6 more variables: arr_delay <dbl>, flight <dbl>, air_time <dbl>,
    distance <dbl>, hour <dbl>, minute <dbl>
data |>
 dplyr::summarise(across(.cols=where(is.numeric), .fns=mean, na.rm=T))
# A tibble: 1 x 14
  year month
                day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <dbl> <dbl> <dbl>
                       <dbl>
                                      <dbl>
                                                 <dbl>
                                                          <dbl>
                                                                         <db1>
1 2013 6.55 15.7
                       1349.
                                      1344.
                                                  12.6
                                                          1502.
                                                                         1536.
# i 6 more variables: arr_delay <dbl>, flight <dbl>, air_time <dbl>,
    distance <dbl>, hour <dbl>, minute <dbl>
```

If applied to a data.frame, summary(), produces a summary of each column. The summary depends on the column type. The output of summary is a shortened version the list of outputs obtained from applying summary to each column (lapply(data, summary)).

```
data |>
  summary()
```

```
day
                                                              sched_dep_time
                  month
                                                  dep_time
     year
Min.
       :2013
              Min. : 1.000
                               Min. : 1.00
                                               Min. : 1
                                                                     : 106
                                                              Min.
1st Qu.:2013
              1st Qu.: 4.000
                               1st Qu.: 8.00
                                               1st Qu.: 907
                                                              1st Qu.: 906
Median:2013
              Median : 7.000
                               Median :16.00
                                               Median:1401
                                                              Median:1359
```

```
: 6.549
Mean
       :2013
               Mean
                                Mean
                                       :15.71
                                                Mean
                                                       :1349
                                                               Mean
                                                                       :1344
3rd Qu.:2013
               3rd Qu.:10.000
                                3rd Qu.:23.00
                                                3rd Qu.:1744
                                                               3rd Qu.:1729
Max.
       :2013
               Max.
                      :12.000
                                Max.
                                       :31.00
                                                Max.
                                                        :2400
                                                               Max.
                                                                       :2359
                                                NA's
                                                       :8255
  dep_delay
                     arr_time
                                 sched_arr_time
                                                  arr_delay
       : -43.00
                                 Min.
                                      : 1
                                                       : -86.000
Min.
                  Min.
                        :
                                                Min.
1st Qu.: -5.00
                  1st Qu.:1104
                                 1st Qu.:1124
                                                1st Qu.: -17.000
Median: -2.00
                  Median:1535
                                 Median:1556
                                                Median : -5.000
Mean
      : 12.64
                  Mean
                        :1502
                                 Mean :1536
                                                Mean
                                                     :
                                                           6.895
3rd Qu.: 11.00
                  3rd Qu.:1940
                                 3rd Qu.:1945
                                                3rd Qu.:
                                                          14.000
       :1301.00
                                        :2359
Max.
                  Max.
                         :2400
                                 Max.
                                                Max.
                                                       :1272.000
NA's
       :8255
                  NA's
                                                NA's
                                                       :9430
                         :8713
  carrier
                       flight
                                    tailnum
                                                        origin
Length: 336776
                   Min.
                        : 1
                                  Length:336776
                                                     Length: 336776
Class : character
                   1st Qu.: 553
                                  Class : character
                                                     Class : character
Mode :character
                   Median:1496
                                  Mode :character
                                                     Mode :character
                          :1972
                   Mean
                   3rd Qu.:3465
                          :8500
                   Max.
                      air_time
                                      distance
                                                       hour
    dest
Length: 336776
                   Min.
                          : 20.0
                                   Min.
                                         : 17
                                                  Min.
                                                         : 1.00
Class :character
                   1st Qu.: 82.0
                                   1st Qu.: 502
                                                  1st Qu.: 9.00
Mode :character
                   Median :129.0
                                   Median: 872
                                                  Median :13.00
                   Mean
                          :150.7
                                         :1040
                                                         :13.18
                                   Mean
                                                  Mean
                   3rd Qu.:192.0
                                   3rd Qu.:1389
                                                  3rd Qu.:17.00
                   Max.
                          :695.0
                                   Max.
                                          :4983
                                                  Max.
                                                         :23.00
                   NA's
                          :9430
                  time_hour
    minute
      : 0.00
                       :2013-01-01 05:00:00.00
Min.
                Min.
1st Qu.: 8.00
                1st Qu.:2013-04-04 13:00:00.00
Median :29.00
                Median :2013-07-03 10:00:00.00
Mean
       :26.23
                       :2013-07-03 05:22:54.64
3rd Qu.:44.00
                3rd Qu.:2013-10-01 07:00:00.00
Max.
       :59.00
                Max.
                       :2013-12-31 23:00:00.00
```

Handling NAs

We add now a few NAs to the data....

```
data2 <- data
data2$arr_time[1:10] <- NA</pre>
```

♦ Houston, we have a problem!

Question

How should we compute the column means now?

```
data2 |>
 dplyr::summarise(across(is.numeric, mean))
# A tibble: 1 x 14
               day dep_time sched_dep_time dep_delay arr_time sched_arr_time
   year month
                                                <dbl>
  <dbl> <dbl> <dbl>
                       <dbl>
                                      <dbl>
                                                         <dbl>
                                                                         <dbl>
1 2013 6.55 15.7
                                      1344.
                                                                         1536.
                          NA
                                                   NA
                                                             NA
# i 6 more variables: arr_delay <dbl>, flight <dbl>, air_time <dbl>,
    distance <dbl>, hour <dbl>, minute <dbl>
```

It is time to look at optional arguments of function mean.

i Question

Decide to ignore NA and to compute the mean with the available data

Solution

```
data2 |>
  dplyr::summarise(across(is.numeric, mean, na.rm=TRUE))
# A tibble: 1 x 14
   year month
                day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <dbl> <dbl> <dbl>
                       <dbl>
                                      <dbl>
                                                 <dbl>
                                                          <dbl>
                                                                         <dbl>
1 2013 6.55 15.7
                                                                         1536.
                       1349.
                                      1344.
                                                  12.6
                                                          1502.
# i 6 more variables: arr_delay <dbl>, flight <dbl>, air_time <dbl>,
    distance <dbl>, hour <dbl>, minute <dbl>
```

i Question

It is possible to remove all rows that contain at least one NA. Show this leads to a different result.

Solution

```
data2 |>
  drop_na() |>
  dplyr::summarise(across(where(is.numeric), mean, na.rm=FALSE)) |>
  knitr::kable()
```

<u>year monthday dep_timbed_depp_timbedary_tismbed_azur_timbediag</u>ht_air_timbistanbour_minute_ 20136.564973.7412748.81340.35912.5555901.92932.81 6.89547943.1250.686348.373.141224.23435

i Question

Compute the minimum, the median, the mean and the maximum of numerical columns

```
data2 |>
 dplyr::select_if(is.numeric) |>
 lapply(function(x) c(med=median(x, na.rm=TRUE),
                     avg=mean(x, na.rm=TRUE),
                     max=max(x, na.rm=TRUE))
$year
med avg max
2013 2013 2013
$month
    med avg max
7.00000 6.54851 12.00000
$day
    med avg max
16.00000 15.71079 31.00000
$dep_time
          avg max
   med
1401.00 1349.11 2400.00
$sched_dep_time
    {\tt med} \qquad {\tt avg} \qquad {\tt max}
1359.000 1344.255 2359.000
$dep_delay
      med
               avg
 -2.00000 12.63907 1301.00000
$arr_time
    med avg max
1536.000 1502.075 2400.000
$sched_arr_time
               max
   med avg
1556.00 1536.38 2359.00
$arr_delay
       \verb| med | avg | max|
 -5.000000 6.895377 1272.000000
$flight
    med avg max
1496.000 1971.924 8500.000
$air_time
            avg max
129.0000 150.6865 695.0000
$distance
```

```
med
              avg
                       max
872.000 1039.913 4983.000
$hour
    med
              avg
                       max
13.00000 13.18025 23.00000
$minute
   med
            avg
                    max
29.0000 26.2301 59.0000
 data2 |>
 dplyr::summarise(across(where(is.numeric),
                          c(median=median, mean=mean, max=max),
                          na.rm=T))
# A tibble: 1 x 42
  year_median year_mean year_max month_median month_mean month_max day_median
                                                                          <dbl>
        <dbl>
                  <dbl>
                           <int>
                                        <dbl>
                                                    <dbl>
                                                              <int>
         2013
                   2013
                            2013
                                                     6.55
1
                                                                 12
                                                                            16
# i 35 more variables: day_mean <dbl>, day_max <int>, dep_time_median <int>,
    dep_time_mean <dbl>, dep_time_max <int>, sched_dep_time_median <dbl>,
    sched_dep_time_mean <dbl>, sched_dep_time_max <int>,
#
#
    dep_delay_median <dbl>, dep_delay_mean <dbl>, dep_delay_max <dbl>,
    arr_time_median <int>, arr_time_mean <dbl>, arr_time_max <int>,
#
#
    sched_arr_time_median <dbl>, sched_arr_time_mean <dbl>,
    sched_arr_time_max <int>, arr_delay_median <dbl>, arr_delay_mean <dbl>, ...
```

Obtain a *nicer* output!

Check with https://dplyr.tidyverse.org/reference/scoped.html?q=funs#arguments

```
Solution
```

```
data2 |>
  dplyr::summarise(across(where(is.numeric),
                          list(median=median,
                               mean=mean,
                               max=max),
                          na.rm=TRUE)
 ) |>
 tidyr::pivot_longer(cols=everything(), names_to="stat", values_to="value") |>
 head() |>
 gt::gt() |>
 gt::fmt_number(columns="value", decimals=2)
                                            value
                        stat
                        year_median
                                         2,013.00
                                         2,013.00
                        year_mean
                        year_max
                                         2,013.00
                        month median
                                            7.00
```

$month_mean$	6.55
$month_max$	12.00

Mimic summary on numeric columns

```
mysum <- data2 |>
 dplyr::summarise(across(is.numeric,
                          list(median=median,
                               mean=mean,
                               max=max,
                               min=min,
                               sd=sd,
                               IQR=IQR) ,
                      na.rm=TRUE))
mysum |>
 tidyr::pivot_longer(
   cols=everything(),
   names_pattern = (\w+)_([a-z]+),
   names_to= c("variable", "stat"),
   values_to="value") |>
 head(20) |>
 gt::gt() |>
 gt::fmt_number(columns="value", decimals=2)
```

stat	value
median	2,013.00
mean	2,013.00
max	2,013.00
\min	2,013.00
sd	0.00
NA	0.00
median	7.00
mean	6.55
max	12.00
\min	1.00
sd	3.41
NA	6.00
median	16.00
mean	15.71
max	31.00
\min	1.00
sd	8.77
NA	15.00
median	1,401.00
	median mean max min sd NA median mean max min sd NA median mean sd NA median median mean max min sd NA

1,349.11 dep_time mean

i Question

Compute a new itinerary column concatenating the origin and dest one. Have a look at Section Operate on a selection of variables

```
Solution
data |>
  dplyr::mutate(itinerary=paste(dest, origin, sep="-")) |>
  dplyr::select(itinerary, dest, origin, everything())
# A tibble: 336,776 x 20
   itinerary dest origin year month
                                          day dep_time sched_dep_time dep_delay
   <chr>
             <chr> <chr> <int> <int> <int> <int>
                                                 <int>
                                                                 <int>
                                                                            <dbl>
 1 IAH-EWR
             IAH
                    EWR
                            2013
                                      1
                                            1
                                                   517
                                                                   515
                                                                                2
2 IAH-LGA
             IAH
                   LGA
                            2013
                                      1
                                            1
                                                   533
                                                                   529
                                                                                4
3 MIA-JFK
             MIA
                    JFK
                            2013
                                            1
                                                                   540
                                                                                2
                                      1
                                                   542
4 BQN-JFK
             BQN
                                      1
                                            1
                    JFK
                            2013
                                                   544
                                                                   545
                                                                               -1
5 ATL-LGA
                                                                               -6
             ATL
                   LGA
                            2013
                                      1
                                            1
                                                   554
                                                                   600
6 ORD-EWR
             ORD
                   EWR
                            2013
                                      1
                                            1
                                                   554
                                                                   558
                                                                               -4
 7 FLL-EWR
                   EWR
                                                                   600
                                                                               -5
             FLL
                            2013
                                      1
                                            1
                                                   555
8 IAD-LGA
             IAD
                    LGA
                            2013
                                      1
                                            1
                                                   557
                                                                   600
                                                                               -3
                                                                               -3
9 MCO-JFK
             MCO
                    JFK
                            2013
                                      1
                                            1
                                                   557
                                                                   600
                                                                               -2
10 ORD-LGA
             ORD
                    LGA
                            2013
                                            1
                                                   558
                                                                   600
                                      1
# i 336,766 more rows
# i 11 more variables: arr_time <int>, sched_arr_time <int>, arr_delay <dbl>,
    carrier <chr>, flight <int>, tailnum <chr>, air_time <dbl>, distance <dbl>,
    hour <dbl>, minute <dbl>, time hour <dttm>
```

Question

3 LGA-EWR

4 SAT-JFK

NA

0.0779

Compute the coefficient of variation (ratio between the standard deviation and the mean) for each itinerary. Can you find several ways?

```
Solution
data |>
  dplyr::mutate(itinerary=paste(dest, origin, sep="-")) |>
  dplyr::select(itinerary, dest, origin, everything()) |>
  dplyr::group_by(itinerary) |>
  dplyr::summarise(coef_var=sd(air_time, na.rm=T)/mean(air_time, na.rm=T), .groups = "dr
  slice_sample(n=10)
# A tibble: 10 x 2
   itinerary coef_var
                <dbl>
   <chr>
 1 SJC-JFK
               0.0475
               0.0829
 2 STL-LGA
```

```
5 DFW-EWR 0.0879
6 SAT-EWR 0.0828
7 STT-JFK 0.0532
8 SEA-JFK 0.0465
9 SAN-JFK 0.0589
10 ATL-LGA 0.0850
```

Compute for each flight the ratio between the distance and the air_time in different ways and compare the execution time (use Sys.time()).

```
Solution
before <- Sys.time()</pre>
data |>
 dplyr::mutate(itinerary=paste(dest, origin, sep="-")) |>
 dplyr::group_by(itinerary) |>
 dplyr::summarize(ratio=mean(air_time)/max(distance)) |>
 dplyr::arrange(desc(ratio))
# A tibble: 224 x 2
   itinerary ratio
   <chr>
            <dbl>
 1 BWI-LGA
            0.219
 2 MEM-JFK 0.172
 3 MYR-LGA 0.166
 4 CAE-LGA 0.164
 5 AVL-LGA 0.154
 6 LEX-LGA
           0.149
 7 SBN-LGA
           0.149
 8 SBN-EWR
            0.147
 9 JAC-JFK
            0.145
10 STL-JFK
            0.145
# i 214 more rows
required_time <- Sys.time() - before</pre>
required_time
Time difference of 0.1288409 secs
```

Question

Which carrier suffers the most delay?

```
data |>
  dplyr::select(carrier, arr_delay) |>
 dplyr::filter(arr_delay > 0) |>
 dplyr::group_by(carrier) |>
 dplyr::summarise(ndelays= n()) |>
# dplyr::arrange(desc(ndelays)) |>
# head(3)
 dplyr::top_n(3, ndelays)
# A tibble: 3 \times 2
  carrier ndelays
  <chr>
          <int>
            23609
1 B6
2 EV
          24484
3 UA
            22222
```

```
Puzzle
year <- 2012L
data |>
  dplyr::select(year, dest, origin) |>
 head()
# A tibble: 6 x 3
   year dest origin
  <int> <chr> <chr>
1 2013 IAH EWR
2 2013 IAH
             LGA
              JFK
3 2013 MIA
4 2013 BQN
              JFK
5 2013 ATL
              LGA
6 2013 ORD
             EWR
data |>
 dplyr::filter(year==year) |>
 dplyr::summarize(n())
# A tibble: 1 x 1
   `n()`
   <int>
1 336776
data |>
  dplyr::filter(year==2012L) |>
 dplyr::summarize(n())
# A tibble: 1 x 1
  `n()`
  <int>
data |>
 dplyr::filter(year==.env$year) |>
 dplyr::summarize(n())
```

```
# A tibble: 1 x 1
   `n()`
   <int>
1     0

data |>
    dplyr::filter(year==.data$year) |>
    dplyr::summarize(n())

# A tibble: 1 x 1
   `n()`
     <int>
1     336776
```

i Question

Can you explain what happens?

Solution

When dplyr::filter(year==year) does year refer to the column of data or to the variable in the global environment?

Flow control

R offers the usual flow control constructs:

- branching/alternative if (...) {...} else {...}
- iterations (while/for) while (...) $\{...\}$ for (it in iterable) $\{...\}$
- function calling callable(...) (how do we pass arguments? how do we rely on defaults?)

If () then {} else {}

If expressions yes_expr and no_expr are complicated it makes sense to use the if (...) {...} else {...} construct

There is also a conditional statement with an optional $else\ \{\}$

```
#| label: if-else
#| eval: false
#| collapse: false
if (condition) {
    ...
} else {
    ...
}
```

i Question

Is there an elif construct in R?

Nope!

R also offers a switch

```
#| label: switch
switch (object,
   case1 = {action1},
   case2 = {action2},
   ...
)
```

```
i There exists a selection function ifelse(test, yes_expr, no_expr).
ifelse(test, yes, no)
Note that ifelse(...) is vectorized.

x <- 1L:6L
y <- rep("odd", 6)
z <- rep("even", 6)

ifelse(x %% 2L, y, z)

[1] "odd" "even" "odd" "even" "odd" "even"

This is a vectorized function</pre>
```

Iterations for (it in iterable) {...}

Have a look at Iteration section in R for Data Science

Question

Create a lower triangular matrix which represents the 5 first lines of the Pascal triangle.

Recall

$$\binom{n}{k} = \binom{n-1}{k-1} + \binom{n-1}{k}$$

5 1 5 10 10 5 1

i Question

Locate the smallest element in a numerical vector

```
Solution
v <- sample(1:100, 100)
v[1:10]
 [1] 11 25 1 17 5 71 8 47 31 21
pmin <- 1
# q: what is the purpose of the following loop?
# for (i in 2:length(v)) {
# if (v[i]<v[pmin]) {</pre>
#
      pmin <- i
#
# }
for (i in seq_along(v)) {
  if (v[i]<v[pmin]) {</pre>
    pmin <- i
  }
}
print(stringr::str_c('minimum is at ', pmin, ', it is equal to ', v[pmin]))
[1] "minimum is at 3, it is equal to 1"
There are some redundant braces {}
```

While (condition) $\{...\}$

i Question

Find the location of the minimum in a vector **v**

```
v <- sample(100, 100)

pmin <- 1  # Minimum in v[1:1]
i <- 2

# q: find le location of the minimum in vector v

while (i <= length(v)) {
    # loop invariant: v[pmin] == min(v[1:i])
    if (v[i]<v[pmin]) {
        pmin <- i
    }
    i <- i + 1
}

print(stringr::str_c('minimum is at ', pmin, ', it is equal to ', v[pmin]))
[1] "minimum is at 52, it is equal to 1"

which.min(v); v[which.min(v)]
[1] 52
[1] 1</pre>
```

Write a loop that checks whether vector \mathbf{v} is non-decreasing.

```
Solution

result <- TRUE

for (i in 2:length(v))
   if (v[i] < v[i-1]) {
     result <- FALSE
     break
   }

if (result) {
   print("non-decreasing")
} else {
   print("not non-decreasing")
}

[1] "not non-decreasing"</pre>
```

Functions

To define a function, whether named or not, you can use the function constructor.

```
foo <- function(arg1, arg2=def2) {
    # body
}</pre>
```

i Question

Write a function that checks whether vector **v** is non-decreasing.

```
Solution
is_non_decreasing <- function(v) {</pre>
  for (i in 2:length(v))
    if (v[i] < v[i-1]) {</pre>
      return(FALSE)
    }
  return(TRUE)
}
is_non_decreasing(v)
[1] FALSE
is_non_decreasing(1:10)
[1] TRUE
A function is an object like any other
is_non_decreasing
function(v) {
  for (i in 2:length(v))
    if (v[i] < v[i-1]) {
      return(FALSE)
    }
  return(TRUE)
}
<bytecode: 0x5a0e391721e8>
body(is_non_decreasing)
{
    for (i in 2:length(v)) if (v[i] < v[i-1]) {
        return(FALSE)
    }
    return(TRUE)
}
args(is_non_decreasing)
function (v)
NULL
```

i Question

Write a function with integer parameter n, that returns the Pascal Triangle with n+1 rows.

```
triangle_pascal <- function(n) {</pre>
 m <- n+1
 T \leftarrow matrix(c(rep(1, m), rep(0, m*(m-1))), nrow=m, ncol=m)
 for (i in 2:m)
   T[i, 2:i] \leftarrow T[i-1, 1:i-1] + T[i-1, 2:i]
 for (i in 1:(m-1))
   T[i, (i+1):m] <- NA
 colnames(T) <- 0:n</pre>
 rownames(T) <- 0:n
 T
}
print(triangle_pascal(10), na.print=" " )
  0 1 2 3 4 5 6 7 8 9 10
0
  1
1
  1 1
2 1 2 1
3 1 3 3
           1
4 1 4 6 4 1
5 1 5 10 10 5 1
6 1 6 15 20 15 6 1
7 1 7 21 35 35 21
                      7
                          1
8 1 8 28 56 70 56 28
                          8 1
9 1 9 36 84 126 126 84 36 9 1
10 1 10 45 120 210 252 210 120 45 10 1
Sanity check: R provides us with function choose
n <- 5
map(0:n, ~ choose(., 0:.))
[[1]]
[1] 1
[[2]]
[1] 1 1
[[3]]
[1] 1 2 1
[[4]]
[1] 1 3 3 1
[[5]]
[1] 1 4 6 4 1
[[6]]
[1] 1 5 10 10 5 1
```

```
t10 <- triangle_pascal(10)

for (n in 0:10)
  for (p in 0:n)
    stopifnot(t10[as.character(n), as.character(p)] == choose(n, p))</pre>
```

How would you generate a Fibonacci sequence of length n?

Recall the Fibonacci sequence is defined by

$$F_{n+2} = F_{n+1} + F_n \qquad F_1 = F_2 = 1$$

```
Solution

fibo <- function(n) {
   res <- integer(n)
   res[1:2] <- 1
   for (k in 3:n) {
      res[k] <- res[k-1] + res[k-2]
   }
   return(res)
}</pre>
fibo(5)

[1] 1 1 2 3 5
```

Question

Write a function that perform binary search in a non-decreasing vector.

```
binary_search <- function(v, x) {
    n <- length(v)
    i <- 1
    j <- n
    while (i <= j) {
        k <- (i+j) %/% 2
        if (v[k] == x) {
            return(k)
        } else if (v[k] < x) {
            i <- k + 1
        } else {
            j <- k - 1
        }
    }
    return(NA)
}</pre>
```

The provided code defines a function binary_search in R, which implements the binary search algorithm to find the position of a target value x within a sorted vector v. The binary search algorithm is efficient, operating in O(log n) time complexity, making it suitable for large datasets.

The function begins by determining the length of the vector v and storing it in the variable n. It then initializes two variables, i and j, to represent the lower and upper bounds of the search range, respectively. Initially, i is set to 1 (the first index of the vector), and j is set to n (the last index of the vector).

The core of the function is a while loop that continues as long as i is less than or equal to j. Within the loop, the midpoint k of the current search range is calculated using integer division (%/%) to ensure it is an integer. The value at index k in the vector v is then compared to the target value x. If the value at v[k] is equal to x, the function returns k, indicating that the target value has been found at this position. If v[k] is less than x, the search range is adjusted by setting i to k+1, effectively discarding the lower half of the current range. Conversely, if v[k] is greater than x, the search range is adjusted by setting j to k-1, discarding the upper half of the current range.

If the loop completes without finding the target value, the function returns NA, indicating that the target value x is not present in the vector v. This implementation assumes that the input vector v is sorted in ascending order.

Code and explanation provided by https://www.techiedelight.com/binary-search-in-r/ through copilot



Read Chapter on functions in Advanced R

In R, argument evaluation is surprising, powerful but taming argument evaluation is real work.

Functional programming

In R, functions are first class entities, they can be defined at run-time, they can be used as function arguments. You can define list of functions, and iterate over them.

Try to use https://purrr.tidyverse.org.

Anonymous functions \(x) body is a shorthand for function (x) { body }

Operators purrr::map_???

i Question

```
Write truth tables for &, |, &&, ||, ! and xor Hint: use purrr::map, function outer()
```

Solution

```
vals <- c(TRUE, FALSE, NA)</pre>
ops <- c(`&`, `|`, `xor`)
truth <- purrr::map(ops, \(x) outer(vals, vals, x))</pre>
names(truth) <- (ops)</pre>
truth
$`.Primitive("&")`
      [,1] [,2] [,3]
[1,] TRUE FALSE
[2,] FALSE FALSE FALSE
[3,]
       NA FALSE
                    NA
$`.Primitive("|")`
     [,1] [,2] [,3]
[1,] TRUE TRUE TRUE
[2,] TRUE FALSE
                  NA
[3,] TRUE
            NA
                  NA
\pi(x, y) \n{\n (x | y) & !(x & y)\n}
      [,1] [,2] [,3]
[1,] FALSE TRUE
[2,] TRUE FALSE
                   NA
[3,]
        NA
              NA
                   NA
```

Question

Write a function that takes as input a square matrix and returns TRUE if it is lower triangular.

Solution

```
lt <- function(A) {
    n <- nrow(A)
    all(purrr::map_lgl(1:(n-1), \(x) all(0 == A[x, (x+1):n])))
}</pre>
```

i Question

Use map, choose and proper use of pronouns to deliver the n first lines of the Pascal triangle using one line of code.

As far as the total number of operations is concerned, would you recommend this way of computing the Pascal triangle?

```
Solution
n <- 5
tp5 <- matrix(unlist(map(0:n,</pre>
           (x) c(choose(x, 0:x), rep(0L, n-x)))),
       nrow=n+1,
       byrow=T)
rownames(tp5) <- 0:n
colnames(tp5) <- 0:n
tp5
  0 1
       2 3 4 5
0100000
          0 0 0
1 1 1 0
2 1 2 1 0 0 0
3 1 3 3
          1 0 0
4 1 4 6 4 1 0
5 1 5 10 10 5 1
No. Using map and choose, we do not reuse previous computations. The total number
of arithmetic operations is \Omega(n^3), it should be O(n^2).
```

Read Chapter on Functional Programming in Advanced R

Further exploration

This notebook walked you through some aspects of R and its packages. We just saw the tip of the iceberg.

We barely mentioned:

- (Non-standard) Lazy evaluation
- Different flavors of object oriented programming
- Connection with C++: RCpp
- Connection with databases: dbplyr
- Building modeling pipelines: tidymodels
- Concurrency
- Building packages
- Building interactive Apps: Shiny
- Attributes (metadata)
- Formulae formula
- Strings stringi, stringr
- Dates lubridate
- and plenty other things

References

- https://www.statmethods.net/index.html
- https://www.datacamp.com/courses/free-introduction-to-r
- dplyr videos

- ggplot2 video tutorial
- cheatsheets

A Readers who really want to learn R should spend time on

- R for Data Science by Wickham, Çetinkaya-Rundel, and Grolemund.
- Advanced R 2nd Edition by Wickham
- Advanced R Solutions by Grosser and Bumann
- Hands-On Programming with R by Grolemund

Don't go without Base R cheatsheet